

# Introduction to Hydra

# Code of conduct

We are operating under the Carpentries **Code of Conduct**.

If you feel that someone has violated this Code of Conduct, please email [si-hpc@si.edu](mailto:si-hpc@si.edu).

# Introductions

# Intended outcomes

After attending this workshop, we hope users come away with these skills:

- How to successfully log in
- How to submit a job
- What to do if something doesn't work
- How to work responsibly on a shared computing resource

# Hydra (SI/HPC)

## People

- Rebecca Dikow (OCIO Data Science Lab), Vanessa González (NMNH GGI), Matt Kweskin (NMNH LAB), and Mike Trizna (OCIO Data Science Lab) provide support for non-CfA users.
- DJ Ding (OCIO) is the full-time Hydra system administrator.
- Sylvain Korzennik (SAO) is the HPC Analyst and provides support for CfA users (and is responsible for the cluster configuration).

## Getting help

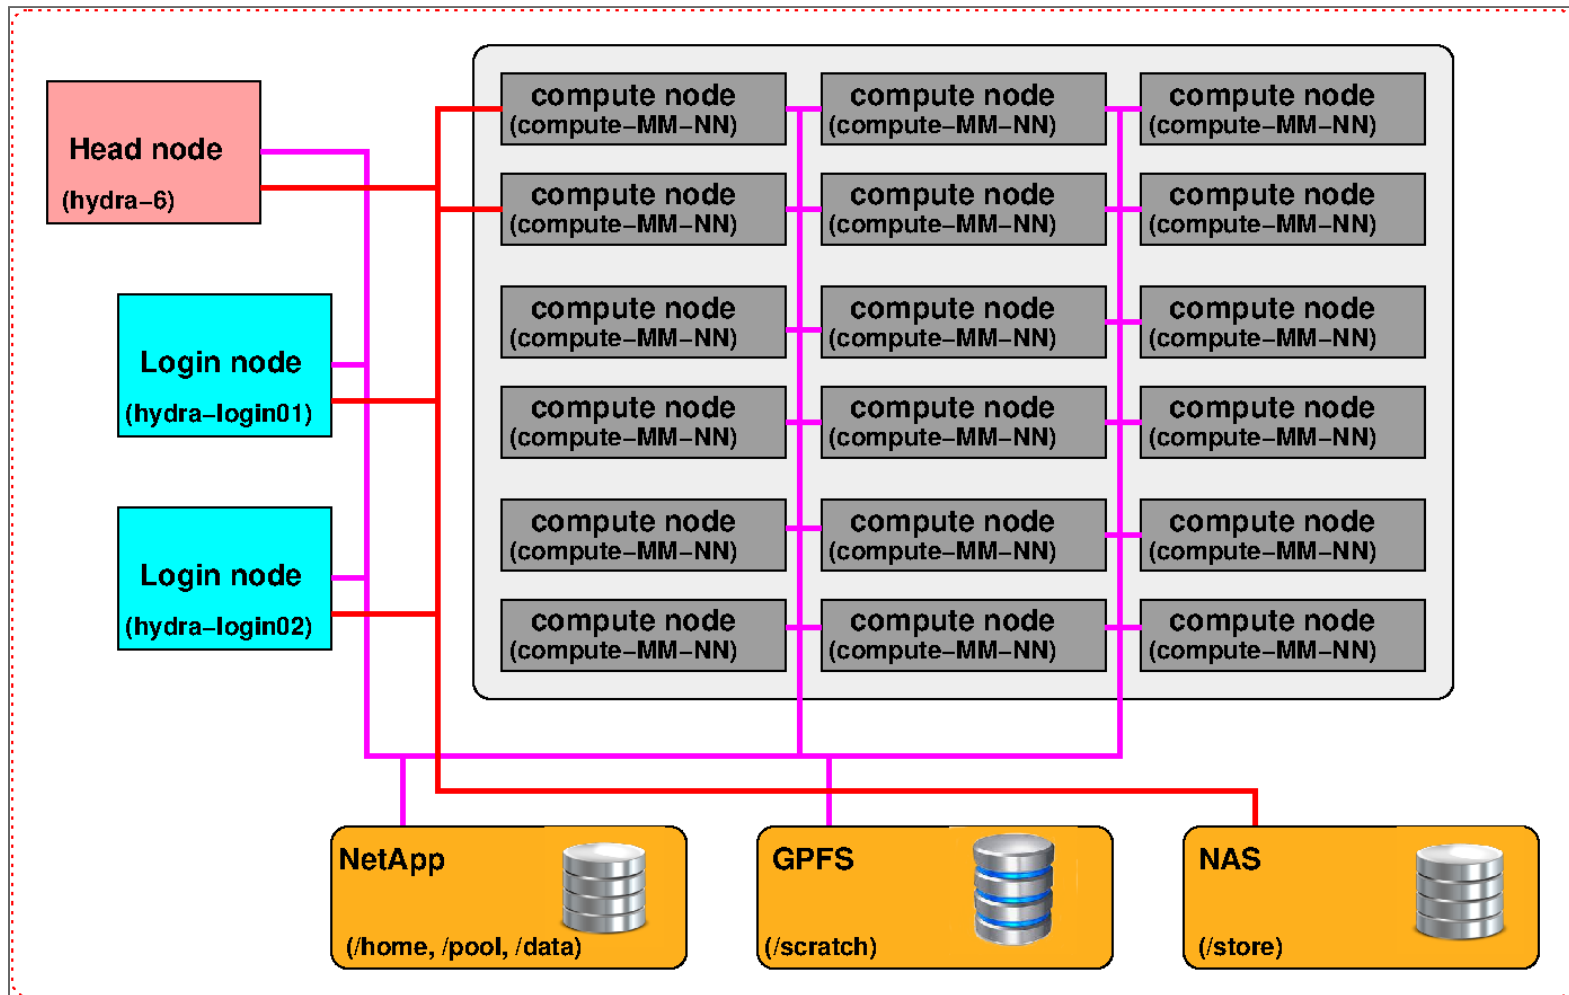
- The Wiki contains detailed documentation
- Email `si-hpc-admin@si.edu` for system-level issues
- For non-CfA users:
  - `Bioinformatics` Brown Bag (Wednesdays, 12-1pm ET, on Zoom)
  - Email `si-hpc@si.edu` (monitored by Rebecca, Vanessa, Matt, and Mike)
- CfA users:
  - email Sylvain or sign up for his office hours.

## Being a good Hydra citizen

- We strive to provide support for users that is inclusive, welcoming, and helps you get your science done.
- We request that users be respectful when asking for help. While we attempt to answer questions rapidly, user support is no one's full-time duties.
- Hydra users agree to abide by its usage policies.



# How is a cluster different than a single-user system?



## How is a cluster different than a single-user system?

- Hydra has 90 compute nodes with between 20 and 128 CPUs each, for a total of 4,896 CPUs.
- Compute nodes have a range of 128GB to 2TB RAM each.
- Two compute nodes have 2 GPUs each.

## Important Takeaways

- Users never need to connect to the Head Node.
- Log in to either `hydra-login01` or `hydra-login02`.
- Do not run commands that use substantial CPU on the login nodes, that's what the compute nodes are for.
- "*Misuse*" of any resources affect the other users (shared resource).

## Disk Storage

- When you log in, you go to your /home directory.
- /home is for your own installed programs and scripts, not for data storage.

## Disk Storage

- Data (large files) belong on `/pool` or `/scratch` and users should run their jobs from here.
- `/pool` and `/scratch` are scrubbed - files older than 180 days are removed.
- `/data` offers some unscrubbed storage for small files (final results, configuration files, etc.)

## Connecting to Hydra

- telework.si.edu (web terminal)
- Mac direct connect (onsite or VPN)
- Windows direct connect (onsite, remote desktop, VPN)
- CfA (VPN, `login.cfa.harvard.edu`, and *trusted* computers: CF/HEA-managed desktops)
- If you don't have an SI VPN but would like to, there is a request form in the SI **ServiceDesk**

## The job scheduler - UGE

- We use UGE (Univa Grid Engine) to schedule resources on Hydra.
- When you submit a job, UGE adds it to the queue and sends it to a compute node with the resources you requested when they become available.

## The job scheduler - UGE

- Each job is assigned a JOB ID, which you can use to check on progress and look at how it used resources when it is complete.
- UGE limits how many resources each user can use concurrently.



## Submitting jobs

- The most common way to run analysis on Hydra is by submitting a job file using the command `qsub`
- We will show you how to build a job file in just a bit
- Users can also start an interactive session using `qrsh`

# Queues

Hydra has different queues to accommodate different resource requests:

- High CPU queues: sThC . q, mThC . q, lThC . q, uThC . q
- High Memory queues: sThM . q, mThM . q, lThM . q, uThM . q

There are other more specialized queues, check the wiki for more information.

# Parallelization

- Depending on the software, you may be able to run a job in **parallel**, which can speed up your analysis.
- Some software uses **threaded** parallelization, where the job is divided across CPUs on a single compute node
- Some software can be compiled to use **MPI** parallelization, where the job is divided across multiple compute nodes

## Parallelization

Look at software documentation to check which kind of parallelization, if any, your software uses and at the [Wiki](#) for how to request the needed resources.

## Parallelization hints

- Some (bioinformatics) software will grab all the CPUs on a compute node unless you tell it otherwise (*not appropriate on a shared machine*).

## Parallelization hints

- Best practice is to use `$NSLOTS` in place of a number of threads in your command. We will demo this in a bit.

# Warnings

- Users that are:
  - Running a job that is inefficient (using <30% of the requested CPU resources), or
  - Running a high-memory job that is using much less than the requested amount of RAM,

will receive an automated warning email.

*We request that you monitor these jobs closely and contact us if you receive repeated warnings*

## Warnings

- When the cluster usage is high, some inefficient jobs will get killed (*fair use* policy).
- Important to receive and read these automatic emails from Hydra (see usage policy).



**Let's Connect**

